

## Lecture 2 tutorial: Application layer

**During tutorials prepare a short report of your activities and show it to your tutor.**

Study the following **questions** and verify the correctness of the **answers** if given.  
Be aware that the exam question might be directly related to the tutorial questions.

**Q1:** HTTP version 2

Read the related part of the HTTP version 2 specifications [RFC 7540](https://tools.ietf.org/html/rfc7540) and describe the structure of the HTTP/2.0 **frame**.

**Q1\_A:** Compile a short answer emphasising the differences between HTTP/1.1

**Q2:** The question below is a bit simplistic, however, if by any chance you have missed something it is a quick way to catch up.

Consider the following string of ASCII characters that were captured by Wireshark when the browser sent an HTTP GET message (i.e., this is the actual content of an HTTP GET message). The characters <cr><lf> are carriage return and line-feed characters (that is, the italicized character string <cr> in the text below represents the single carriage-return character that was contained at that point in the HTTP header). Answer the following questions, indicating where in the HTTP GET message below you find the answer.

```
GET /cs453/index.html HTTP/1.1<cr><Lf>
Host: gaia.cs.umass.edu <cr><lf>
User-Agent: Mozilla/5.0 (Windows;U; Windows NT 5.1; en-US; rv:1.7.2)
Gecko/20040804 Netscape/7.2 (ax) <cr><lf>
Accept: ext/xmL, application/xml, application,/xhtml+xml, text/html;
q=0.9,text/plain; q=g.8,image/png,*/*; q=0.5<cr><lf>
Accept-Language: en-us, en;q=0.5 <cr><lf>
Accept-Encoding: zip, deflate<cr><lf>
Accept-Charset: ISO-8859-1, utf-8;q=0.7,*;q=0.7<cr><lf>
Keep-Alive: 300<cr><lf>
Connection: keep-alive<cr><lf>
<cr><lf>
```

- What is the URL of the document requested by the browser?
- What version of HTTP is the browser running?
- Does the browser request a non-persistent or a persistent connection?
- What is the IP address of the host on which the browser is running?
- What type of browser initiates this message? Why is the browser type needed in an HTTP request message?

**Q2\_A:**

- The document request was `http://gaia.cs.umass.edu/cs453/index.html`. The Host: field indicates the server's name and `/cs453/index.html` indicates the file name.
- The browser is running HTTP version 1.1, as indicated just before the first <cr><lf> pair.

- c) The browser is requesting a persistent connection, as indicated by the Connection: keep-alive.
- d) This is a trick question. This information is not contained in an HTTP message anywhere. So there is no way to tell this from looking at the exchange of HTTP messages alone. One would need information from the IP datagrams (that carried the TCP segment that carried the HTTP GET request) to answer this question.
- e) Mozilla/5.0. The browser type information is needed by the server to send different versions of the same object to different types of browsers.

**Q3:** The text below shows the reply sent from the server in response to the HTTP GET message in the question above. Answer the following questions, indicating where in the message below you find the answer.

- a. Was the server able to successfully find the document or not? What time was the document reply provided?
- b. When the document was last modified?
- c. How many bytes are there in the document being returned?
- d. What are the first 5 bytes of the document being returned? Did the server agree to a persistent connection?

```
HTTP/1.1 200 OK<cr><lf>
Date: Tue, 07 Mar 2008 12:39:45GMT<cr><lf>
Server: Apache/2.0.52 (Fedora)<cr><lf>
Last-Modified: Sat, 10 Dec2005 18:27:46 GMT<cr><lf>
ETag: "u526c3-f22-a88a4c80"<cr><lf>
Accept-Ranges: bytes<cr><lf>
Content-Length: 3874<cr><lf>
Keep-Alive: timeout=max=100<cr><lf>
Connection: Keep-Alive<cr><lf>
Content-Type: text/html; charset=ISO-8859-1<cr><lf>
<cr><lf>
<!doctype html public "-//w3c//dtd html 4.0 transitional//en"><lf>
<html><lf> <head><lf> <meta http-equiv="Content-Type"
content="text/html; charset=iso-8859-1"><lf> <meta name="GENERATOR"
content="Mozilla/4.79 [en] (Windows NT 5.0; U) Netscape]"><lf>
<title>CMPSCI 453 / 591 / NTU-ST550A Spring 2005
homepage</title><lf></head><lf>
<much more document text following here (not shown)>
```

### Q3\_A:

- a) The status code of 200 and the phrase OK indicate that the server was able to locate the document successfully. The reply was provided on Tuesday, 07 Mar 2008 12:39:45 Greenwich Mean Time.
- b) The document index.html was last modified on Saturday 10 Dec 2005 18:27:46 GMT.
- c) There are 3874 bytes in the document being returned.

- d) The first five bytes of the returned document are : <!doc. The server agreed to a persistent connection, as indicated by the Connection: Keep-Alive field

**Q4:** We used the [Rex Swain's HTTP Viewer](#) as a HTTP analyzer. This side is no longer available (re-check !).

Find another HTTP analyser to **request** and get the **response** from the site as above:

[gaia.cs.umass.edu/cs453/index.html](http://gaia.cs.umass.edu/cs453/index.html)

Compare changes in the request/response wrt Q2 and 3

**(Answer the questions as in Q2 and Q3)**

**Q5:** Consider a redirection problem. My web page originally existed on the server [www.csse.monash.edu.au/~app](http://www.csse.monash.edu.au/~app). If you click on that side, you will be redirected to [users.monash.edu/~app/](http://users.monash.edu/~app/). The question is: find the HTTP source code on the page [www.csse.monash.edu.au/~app](http://www.csse.monash.edu.au/~app) that performs the redirection. A nice HTTP analyzer would be handy. If you cannot find one, use the Wireshark.

**Q6:** Suppose within your Web browser you click on a link to obtain a Web page. The IP address for the associated URL is not cached in your local host, so a DNS lookup is necessary to obtain the IP address. Suppose that  $n$  DNS servers are visited before your host receives the IP address from DNS; the successive visits incur an  $RTT_{DNS}$  of  $RTT_1, \dots, RTT_n$ . Further suppose that the Web page associated with the link contains exactly one object, consisting of a small amount of HTML text. Let  $RTT_0$  denote the RTT between the local host and the server containing the object. Assuming zero transmission time of the object, how much time elapses from when the client clicks on the link until the client receives the object?

**Q6\_A:**

The total amount of time to get the IP address is

$$RTT_{DNS} = RTT_1 + \dots + RTT_n$$

Once the IP address is known,  $RTT_0$  elapses to set up the TCP connection and another  $RTT_0$  elapses to request and receive the small object. The total response time is

$$RTT_{DNS} + 2RTT_0$$

**Q7:** Referring to question **Q6**, suppose the HTML file references **eight very small objects** on the same server. Neglecting transmission times, how much time elapses with

- Non-persistent HTTP with no parallel TCP connections?
- Non-persistent HTTP with the browser configured for 5 parallel connections?
- Persistent HTTP?

**Q7\_A:**

a)

$$RTT_{DNS} + 2RTT_0 + 8 \cdot 2RTT_0 = RTT_{DNS} + 18RTT_0$$

( $2RTT_0$  one to establish TCP connection, the other to send HTTP request – response)

b)


$$RTT_{DNS} + 2RTT_0 + 2 \cdot 2RTT_0 = RTT_{DNS} + 6RTT_0$$

c)

$$RTT_{DNS} + 2RTT_0 + RTT_0 = RTT_{DNS} + 3RTT_0$$

**Q8:** Refer to lecture slides: **27 – 30**

An institutional network is connected to the Internet in the following way:



Suppose that the average object size is 850,000 bits and that the average request rate from the institution's browsers to the origin servers is 16 requests per second. Also suppose that the amount of time it takes from when the router on the Internet side of the access link forwards an HTTP request until it receives the response is three seconds on average.

Model the total average response time as the sum of the average access delay (that is, the delay from Internet router to institution router), and the average Internet delay. For the average access delay, use  $\Delta/(1 - \Delta\beta)$ , where  $\Delta$  is the average time

required to send an object over the access link and  $\beta$  is the arrival rate of objects to the access link.

- a. Find the total average response time.
- b. Now suppose a cache is installed in the institutional LAN. Suppose the miss rate is 0.4. Find the total response time.

**Q8\_A:**

- a) The time to transmit an object of size  $L$  over a link of rate  $R$  is  $L/R$ . The average time is the average size of the object divided by  $R$ :

$$\Delta = (850,000 \text{ bits}) / (15,000,000 \text{ bits/sec}) = .0567 \text{ sec}$$

The traffic intensity on the link is given by

$$\beta\Delta = (16 \text{ requests/sec})(.0567 \text{ sec/request}) = 0.907.$$

Thus, the average access delay is  $(.0567 \text{ sec}) / (1 - .907) \approx .6 \text{ seconds}$ .

The total average response time is therefore  $.6 \text{ sec} + 3 \text{ sec} = 3.6 \text{ sec}$ .

- b) The traffic intensity on the access link is reduced by 60% since the 60% of the requests are satisfied within the institutional network.

Thus the average access delay is

$$(.0567 \text{ sec}) / [1 - (.4)(.907)] = .089 \text{ seconds}.$$

The response time is approximately zero if the request is satisfied by the cache (which happens with probability .6);

The average response time is

$$.089 \text{ sec} + 3 \text{ sec} = 3.089 \text{ sec}$$

for cache misses (which happens 40% of the time).

So the average response time is

$$(.6)(0 \text{ sec}) + (.4)(3.089 \text{ sec}) = 1.24 \text{ seconds}$$

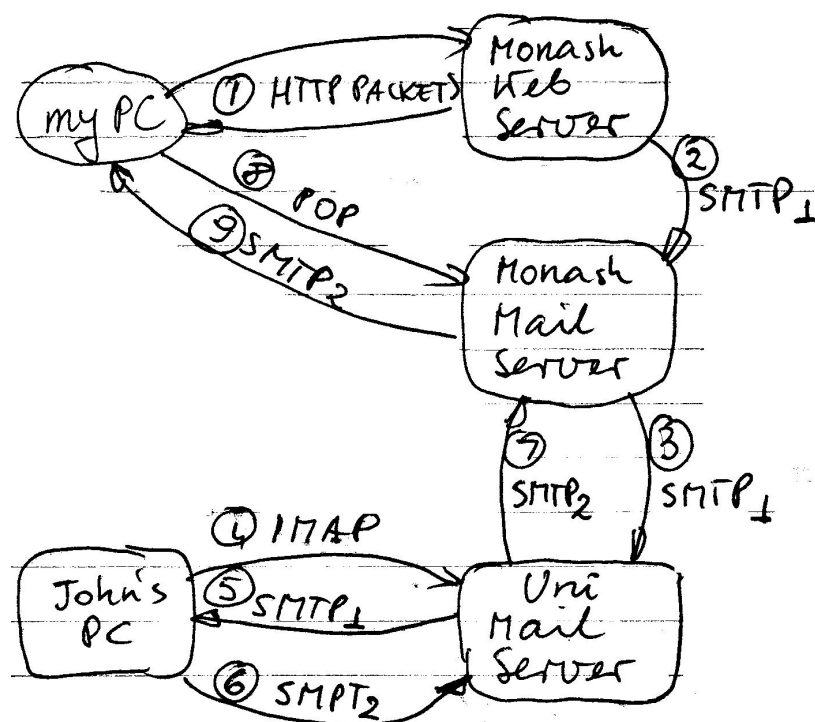
Thus the average response time is reduced from 3.6 sec to 1.24 sec.

**Q9: E-mail. (a good exam question)**

Assume that:

- You use a **web email** at **Monash** to send an email to [John@uni.edu](mailto:John@uni.edu).
  - John is using an **IMAP mailer** to check his e-mail and respond to your e-mail.
  - You use a **POP mailer** to check John's response.
- a. Illustrate in a **single drawing** the above scenario drawing the flow of packets between the computers involved. Annotate all servers and client computers involved.
  - b. Give a brief step-by-step explanation of the flow of packets. Number the steps and related packets in the drawing.

Q9\_A:



1. I invoke a web-mailer and exchange HTTP packets with the Web server to prepare e-mail to be sent.
2. When the e-mail is completed the Monash Web server sends the SMTP-1 packet to the Monash mail server
3. The Monash mail server sends the SMTP-1 packet to the Uni mail server. My e-mail is saved in John's mailbox.
4. John reads my e-mail from his mailbox using an IMAP mailer. Relevant IMAP request is sent to the Uni mail server.
5. In response, the Uni mail server sends the SMTP-1 packet to John's PS.
6. To reply to my e-mail John invokes his IMAP mailer and the resulting SMTP-2 packet is sent to the Uni mail server.
7. The Uni mail server sends the SMTP-2 packet to the Monash mail server. John's reply is saved in my mailbox.
8. I read John's reply using a POP mailer. A relevant POP packet is sent to the Monash mail server
9. The Monash mail server sends to me the SMTP-2 packet containing John's reply

**Q10:** Refer to the slide **36** and a relevant **SMTP** standard and explain how the Message-ID is created. Show where in the standard you have found the relevant information

**Q10\_A:** Please write a good answer.

**Q11:** Briefly describe the **MIME** standard. Note many related RFCs